

Brief History of SQL

- ❑ Developed 1974 at IBM as the language of the relational DBMS System R
- ❑ Original name: SEQUEL = Structured English Query Language
- ❑ Versions
 - ❖ SQL-86 (1986/87): First formalization by ANSI and first ISO standard
 - ❖ SQL-89 (1989): Addition of integrity constraints
 - ❖ SQL-92 (1992): Also denoted as SQL2, major revision
 - ❖ SQL:1999 (1999): Also denoted as SQL3, addition of object-oriented features, triggers, regular expressions
 - ❖ SQL:2003: Addition of XML-related features, standardized sequences
 - ❖ SQL:2006: Addition of features to import and store XML data in an SQL database
 - ❖ SQL:2008: Minor revision
 - ❖ SQL:2011: Addition of support for temporal databases
 - ❖ SQL:2016: Addition of row pattern matching and JSON support
- ❑ Database vendors provide special extensions (packages) and SQL dialects

Components of SQL (I)

❑ Data definition language (DDL)

- ❖ Commands for creating and changing the data structures for the three levels of a database (external levels, conceptual level, physical level) such as the
 - definition of relation schemas
 - deletion of relation schemas
 - creation of index structures
 - modification of relation schemas
- ❖ Remember: *Relations (tables)* are based on a *set (list)* model

❑ Data manipulation language (DML)

- ❖ Update commands for
 - inserting data objects (**tuples**)
 - modifying tuples
 - deleting tuples
- ❖ Interactive formulation of **queries**

Components of SQL (II)

❑ Embedded SQL and dynamic SQL

- ❖ Mechanisms to embed SQL statements into general-purpose programming languages
- ❖ Examples of such **host languages** are Fortran, C, C++, or Java

❑ Integrity

- ❖ Commands for specifying integrity constraints that the data stored in the database must satisfy
- ❖ Attempts to violate these integrity constraints are disallowed and blocked

❑ View definition

- ❖ Commands for defining views
- ❖ Interesting aspect: Views are defined by an SQL query

Components of SQL (III)

❑ Transaction control

- ❖ Explicit commands for specifying the beginning and ending of transactions
- ❖ Transactions are usually started and ended in the background
- ❖ Explicit aborting a transaction is also possible

❑ Authorization

- ❖ Commands for specifying access rights to relations and views
- ❖ Different user groups usually obtain different access rights

SQL Data Types (I)

- ❑ SQL data types are part of the SQL DDL, built-in types, and serve as **domains**, that is, **attribute data types**
- ❑ **char(n), character(n)**
 - ❖ Character strings of fixed length n
 - ❖ Length is user-specified
 - ❖ All character strings have the same length, strings are padded with blank characters to the right if needed
- ❑ **varchar(n) , char varying(n), character varying(n)**
 - ❖ Character strings of variable maximum length n
 - ❖ Length is user-specified
 - ❖ Less than n characters are stored in shorter strings
- ❑ **int, integer**
 - ❖ Data type for a finite subset of the integers
 - ❖ Machine-dependent (e.g., 4 bytes)

SQL Data Types (II)

❑ **smallint**

- ❖ Small integers
- ❖ Machine-dependent subset of the integer type (e.g., 2 bytes)

❑ **numeric(p , d), decimal(p , d), dec(p , d)**

- ❖ Fixed-point numbers with user specified precision
- ❖ Exact decimal numbers
- ❖ p = total number of digits (plus a sign), d = number of the p digits to the right of the decimal point
- ❖ Example: **numeric**(3,1) allows 87.3 to be stored exactly, this is not the case for 765.4 or 0.19

❑ **real**

- ❖ Floating-point numbers with machine-dependent precision

❑ **double precision**

- ❖ Double-precision floating-point numbers with machine-dependent precision

SQL Data Types (III)

❑ **boolean**

- ❖ Has the traditional values of *true* and *false*
- ❖ Because of the presence of *null* values, a three-valued logic is used with the additional possible Boolean value *unknown*
- ❖ Any comparison involving the *null* value or an *unknown* truth value returns an *unknown* result

❑ **date**

- ❖ Calendar date with year (4 digits), month (2 digits), and day (2 digits)
- ❖ Format: YYYY-MM-DD
- ❖ Example: **date** '2015-04-13'

❑ **time**

- ❖ The time of day, in hours, minutes, and seconds (2 digits each)
- ❖ Format: HH:MM:SS
- ❖ Example: **time** '08:34:15'
- ❖ Local time zone is assumed

SQL Data Types (IV)

❑ **time with time zone**

- ❖ Time difference to Greenwich Mean Time included
- ❖ Displacement from the standard universal time zone in the range of +13:00 to -12:59 in units of HH:MM

❑ **timestamp**

- ❖ Merges date and time fields
- ❖ Format: YYYY-MM-DD HH:MM:SS
- ❖ Example: **timestamp** '2015-04-13 08:34:15'
- ❖ Optional **with time zone** qualifier

❑ **interval**

- ❖ Intervals are *relative values* that can increment or decrement an absolute value of the types **date**, **time**, or **timestamp**
- ❖ Two classes of intervals
 - Year-month intervals containing the year and month fields
 - Day-time intervals containing a contiguous selection of day, hours, minute, second

SQL Data Types (V)

❑ **bit(*n*)**

- ❖ Bit string of fixed length *n* bits

❑ **bit varying(*n*)**

- ❖ Bit string of varying length with a maximum length of *n* bits

❑ **blob**

- ❖ **Binary large objects**
- ❖ Byte sequences of variable length up to a user-defined limit
- ❖ Examples: image **blob**(10KB), movie **blob**(2GB)
- ❖ Maximum length can be specified in kilobits (K), megabits (M), gigabits (G), kilobytes (KB), megabytes (MB), or gigabytes (GB)
- ❖ Used to store **complex objects** from advanced applications, such as multimedia objects, video clips, spatial objects, high-resolution medical images, satellite images, proteins, enzymes, times series
- ❖ Low-level operations for binary reads and writes only

SQL Data Types (VI)

❑ **clob**

- ❖ Character large objects
- ❖ Large character strings to store books and large documents
- ❖ Example: textbook **clob**(1 GB)

❑ Declaration of a domain

- ❖ Command: **create domain**
- ❖ Example: **create domain** SSN_Type **as char(9)**
- ❖ Advantages
 - ❖ Simple centralized change of a data type for a domain which is used by several attributes in a database schema
 - ❖ Domains improve schema readability
- ❖ Disadvantage: No complex declarations possible

Integrity Constraints and Default Values (I)

- ❑ In the following: The most important commands of the SQL **Data Definition Language (DDL)**
- ❑ Integrity constraints
 - ❖ are conditions that restrict the possible database states
 - ❖ ensure the consistency of a database
 - ❖ will be discussed in much more detail in a later lecture
- ❑ Since SQL allows null values (**null**), an integrity constraint **not null** can be defined, if for a specific attribute a null value is *not* allowed
 - ❖ Example 1: To store a student's UFID without storing the student's name does not make much sense
 - ❖ Example 2: A lecture code (e.g., CIS 4301) should always be accompanied by the lecture title ("Information and Database Management Systems I")
- ❑ It is recommended to specify the **not null** condition explicitly for each primary key, although this holds implicitly already for each primary key

Integrity Constraints and Default Values (II)

□ Default values

- ❖ Definition of a default value for an attribute possible by attaching the clause **default** <value> to the attribute definition
- ❖ A default value is inserted into each new tuple if an explicit value for this attribute is not specified
- ❖ If a default clause is not defined, the default value is *null*

□ Primary keys

- ❖ The clause **primary key** (A_1, \dots, A_n) specifies that the attributes $A_1, \dots, A_n \in \mathcal{R}$ ($n \geq 1$) form the primary key of the relation R with respect to \mathcal{R}
- ❖ The set of selected attributes must be unique and minimal
- ❖ Let $\mathcal{R}(A_1, \dots, A_n, B_1, \dots, B_m)$ be a relation schema, A_1, \dots, A_n be the primary key, B_1, \dots, B_m be nonprime attributes, and R be a relation with schema \mathcal{R} . Then the following holds for all tuples of R :

$$\forall s, t \in R, s \neq t : \neg ((s.A_1, \dots, s.A_n) = (t.A_1, \dots, t.A_n))$$

Integrity Constraints and Default Values (III)

❑ Foreign keys

- ❖ A foreign key value references (“points to”) a primary key value in another relation to represent an $m:1$ -relationship of an E-R diagram
- ❖ Definition of a foreign key by the **foreign key clause**
- ❖ Foreign keys ensure **referential integrity**, that is, there are no inconsistent or dangling references
- ❖ A foreign key can or cannot become (part of) the primary key of the relation schema into which it is imported

assistants					professors			
pers-id	name	room	boss		pers-id	name	rank	room
3002	Platon	156	2125	→	2125	Sokrates	C4	226
3003	Aristoteles	199	2125	→	2126	Russel	C4	232
3004	Wittgenstein	101	2126	→	2127	Kopernikus	C3	310
3005	Rhetikus	130	2127	→	2133	Popper	C3	052
3006	Newton	120	2127	→	2134	Augustinus	C3	309
3007	Spinoza	155	2134	→	2136	Curie	C4	036
					2137	Kant	C4	007

Integrity Constraints and Default Values (IV)

❑ Candidate keys

- ❖ Any unique and minimal set of attributes that could serve but has not been selected as a primary key
- ❖ The fact that the attributes A_1, \dots, A_n form a candidate key is specified by the integrity constraint **unique** (A_1, \dots, A_n)
- ❖ Later it will be an important task to find *all* candidate keys of a relation schema

Creation of a Relation Schema / Table Schema (I)

- ❑ Precisely speaking, there are no relations in SQL but tables (set model versus list model)
- ❑ Creation of a schema with the aid of the DDL clause

create table R (A_1 D_1 , ..., A_n D_n ,

[< integrity constraint₁ >, ..., < integrity constraint_k >])

R relation name, A_i name of an attribute in the schema of relation R , D_i domain of A_i

- ❑ In BNF notation
 - ❖ **create table** <relation name>(<relation comp> [, <relation comp>]*)
 - ❖ <relation comp> ::= <column definition> | <integrity constraint>
 - ❖ <column definition> ::= <attribute name> <type>
[<default value> | **not null** | **unique**]
 - ❖ <default value> ::= [**default** <literal> | **null**]
 - ❖ The exact treatment of integrity constraints is discussed later.

Creation of a Relation Schema / Table Schema (II)

- Example: University schema (with the main integrity constraints)

```
create table students
  (reg-id int not null,
   name  varchar(30) not null,
   sem   int,
   primary key (reg-id));
```

```
create table professors
  (pers-id int not null,
   name    varchar(30) not null,
   room    int unique,
   rank    char(2),
   primary key (pers-id));
```

```
create table assistants
  (pers-id int not null,
   name    varchar(30) not null,
   room    int unique,
   boss    int,
   primary key (pers-id),
   foreign key (boss) references
     professors(pers-id));
```

```
create table lectures
  (id      int not null,
   title   varchar(30),
   credits int,
   held_by int,
   primary key (id),
   foreign key (held_by) references
     professors(pers-id));
```


Creation of a Relation Schema / Table Schema (III)

- Example: University schema (with the main integrity constraints, continued)

```
create table attends
  (reg-id  int not null,
   id      int not null,
   primary key (reg-id, id),
   foreign key (reg-id) references
        students(reg-id),
   foreign key (id) references
        lectures(id));
```

```
create table test
  (reg-id  int not null,
   id      int not null,
   pers-id int not null,
   grade   numeric(2,1),
   primary key (reg-id, id, pers-id),
   foreign key (reg-id) references students(reg-id),
   foreign key (id) references lectures(id),
   foreign key (pers-id) references professors(pers-id));
```


```
create table is_precondition_of
  (predecessor int not null,
   successor   int not null,
   primary key (predecessor,
        successor),
   foreign key (predecessor)
        references lectures(id),
   foreign key (successor)
        references lectures(id));
```

Modification of a Relation Schema / Table Schema


- ❑ Adding an attribute (a new column) to an *existing* relation
 - ❖ **alter table** <relation name> **add** <column definition>
 - ❖ All tuples in the relation are assigned *null* as the value for the new attribute
 - ❖ Constraint **not null** is only allowed if a default value is specified
- ❑ Deleting an attribute (a column) from an *existing* relation
 - ❖ **alter table** <relation name> **drop** <column definition>
 - ❖ Command will be declined if integrity constraints would be violated
 - ❖ Example 1: Delete a column that has values which are referenced by tuples from other relations
 - ❖ Example 2: Delete a column that forms the primary key of a relation
- ❑ Deletion of the schema and instance (that is, data) of a relation
 - ❖ **drop table** <relation name>
- ❑ Deletion of the instance (that is, data) of a relation but not its schema
 - ❖ **delete from** <relation name>

Creation of an Index (I)

- ❑ An **index** is a *persistent* data structure that provides accelerated access to the rows of a table based on the values of one or more attributes
- ❑ Goal: Significant improvement of query response time
- ❑ Example: List of tuples that store first and last names of persons



- ❑ A search for Tom Smith above takes $O(n)$ time (n number of tuples)




- ❑ A search for Tom Smith takes $O(\log n + k)$ time (k number of “Smith” tuples)
- ❑ A measure for the efficiency is, in general, the number of page accesses to the hard disk

Creation of an Index (II)

- ❑ Disadvantage: Indexes have to be updated by the system every time the underlying tables are updated
 - Too many indexes lead to a large overhead and slow down the performance of the DBMS, right balance needed
- ❑ SQL command for creating an index
 - ❖ **create** [**unique**] **index** <index name> **on** <relation name>
(<attribute name> [<order>] [, <attribute name> [<order>]]*) [**cluster**]
 - ❖ <order> ::= **Asc** | **Desc**
 - ❖ **unique** means that two tuples with the same values for all indexed attribute names are forbidden (indexed attributes fulfil key condition)
 - ❖ **cluster** means that the tuples of the relation are actually inserted into the index structure by physical proximity (only one cluster index per relation)
- ❑ Example: **create unique index** room_index **on** professors (room);
- ❑ SQL command for deleting an index
 - ❖ **drop index** <index name>

Creation of Views (I)

- ❑ A **view** is a *virtual relation* that does not exist persistently in the database but can be produced upon request by a particular user, at the time of request



- ❑ Views are regarded as derived relations which are defined by queries
- ❑ **create view** <view name> [(<attribute name> [, <attribute name>]*)]
as <subquery>

Creation of Views (II)

- ❑ Example

```
create view major_students as  
    select * from students where sem > 4;
```

- ❑ The keyword “*” is a shortcut for the complete attribute list of those relations placed after **from**

- ❑ Deletion of views

```
drop view <view name>
```