

## Lossless Join Decomposition (XI)


- ❑ The goal is to use  $F$  to prove that  $t \in r$  holds
- ❑ Strategy of the Chase test
  - ❖ We *chase* the matrix by applying the FDs in  $F$  to equate symbols in the matrix whenever possible
  - ❖ If we manage to obtain a row that is equal to  $t$  (that is, the row only contains unsubscripted letters), we have proved that any tuple  $t$  in the join of the projections was actually a tuple of  $R$
- ❑ Example
  - ❖ We continue our previous example and assume the set  $F = \{A \rightarrow B, B \rightarrow C, CD \rightarrow A\}$  of FDs
  - ❖ The FDs in  $F$  can be applied in any order and several times

$A$	$B$	$C$	$D$
$a$	$b_1$	$c_1$	$d$
$a$	$b_2$	$c$	$d_2$
$a_3$	$b$	$c$	$d$

$\xrightarrow{A \rightarrow B}$

$A$	$B$	$C$	$D$
$a$	$b_1$	$c_1$	$d$
$a$	$(b_1)$	$c$	$d_2$
$a_3$	$b$	$c$	$d$


## Lossless Join Decomposition (XII)



- ❖ The last row has become equal to  $t$
- ❖ We have shown that if  $r$  satisfies the FDs  $A \rightarrow B$ ,  $B \rightarrow C$ , and  $CD \rightarrow A$ , then whenever we project  $r$  onto  $\{A, D\}$ ,  $\{A, C\}$ , and  $\{B, C, D\}$  and rejoin, what we get must have been in  $r$

### □ Example

- ❖ Consider the relation  $R(A, B, C, D)$  with  $F = \{B \rightarrow AD\}$  and the decomposition  $\{A, B\}$ ,  $\{B, C\}$ , and  $\{C, D\}$ . Apply the Chase test.

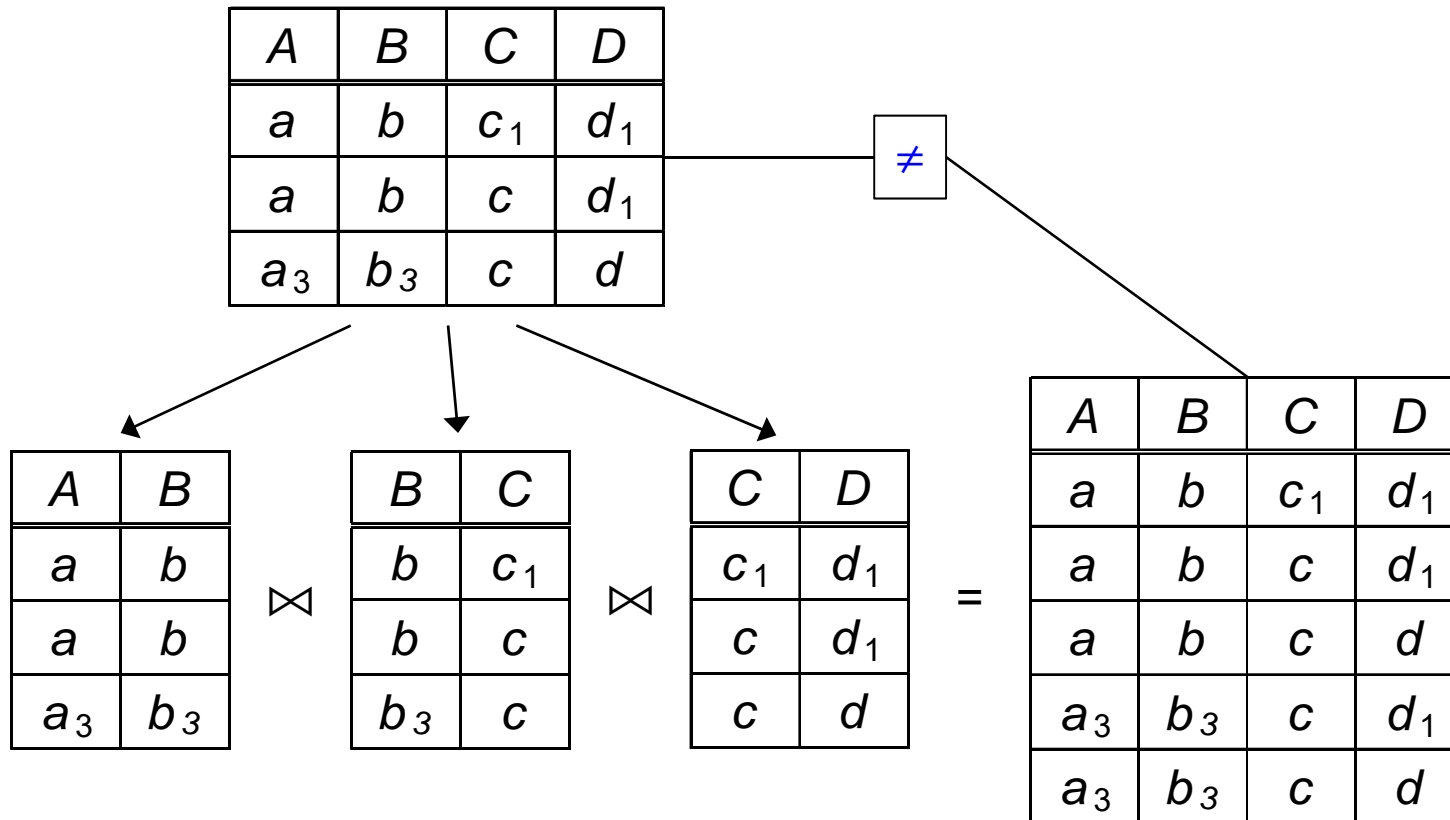


There is no row that is fully unsubscripted. The decomposition is lossy.

## Lossless Join Decomposition (XIII)

### ❑ Example (*continued*)

- ❖ Another way to show this is: Treat the right table as a relation with three tuples, decompose it, and then rejoin



## Dependency Preservation (I)

- ❑ For performance reasons it would be useful if each FD in  $F$  either could be checked directly in one of the relation schemas  $R_i$  of the decomposition, or could be inferred from the FDs that hold on the attributes of some schema  $R_i$
- ❑ If one of the FDs is not represented in some schema  $R_i$  of the decomposition, we cannot check and enforce this constraint on a single relation but have to join multiple relations in order to include all attributes involved in that FD
- ❑ Given a set  $F$  of FDs on a relation schema  $R$  and a decomposition  $R_1, \dots, R_n$  of  $R$ , the **restriction**  $F_{R_i}$  of  $F$  to  $R_i$  is defined as  $F_{R_i} = \{X \rightarrow Y \in F^+ \mid X \cup Y \subseteq R_i\}$
- ❑ All FDs of  $F_{R_i}$  can be checked *locally* on  $R_i$  *alone* (without the need of joins)
- ❑ All FDs in  $F^+$  are used, not only the FDs in  $F$
- ❑ Example
  - ❖ Suppose we have  $R(A, B, C)$ ,  $F = \{A \rightarrow B, B \rightarrow C\}$ ,  $R_1(A, C)$ ,  $R_2(A, B)$
  - ❖  $F_{R_1}$  includes  $A \rightarrow C$  since  $A \rightarrow C \in F^+$  but  $A \rightarrow C \notin F$

## Dependency Preservation (II)

- ❑ Question: Is testing only the restrictions sufficient?
- ❑  $F_r = F_{R_1} \cup F_{R_2} \cup \dots \cup F_{R_n}$  is a set of FDs on  $R$ , i.e.,  $F_r \subseteq F^+$
- ❑ In general,  $F_r \neq F$  holds
- ❑ Even if  $F_r \neq F$  holds, it can be that  $F_r^+ = F^+$  holds, i.e.,  $F_r \equiv F$
- ❑ If  $F_r^+ = F^+$  holds, then every FD in  $F$  is logically implied by  $F_r$
- ❑ A decomposition having the property  $F_r^+ = F^+$  is a **dependency-preserving decomposition**
- ❑ If each FD in  $F$  can be tested on one of the relation schemas of the decomposition, the decomposition is dependency-preserving
- ❑ The following algorithm tests for dependency preservation in general
- ❑ It is expensive since it requires the computation of  $F^+$

## Dependency Preservation (III)

- ❑ Example: Let  $R(A, B, C)$  and  $F = \{A \rightarrow B, B \rightarrow C, C \rightarrow A\}$ . Let  $R_1(A, B)$  and  $R_2(B, C)$  be a decomposition of  $R$ . Check whether the decomposition is dependency-preserving.
- ❑ It is easy to see that  $A \rightarrow B \in F_{R_1}$  and  $B \rightarrow C \in F_{R_2}$
- ❑ The question is whether the decomposition preserves the FD  $C \rightarrow A$
- ❑ Further question: Does the fact that  $A$  and  $C$  are contained together neither in  $R_1$  nor in  $R_2$  mean that the decomposition is not dependency-preserving?
- ❑  $F^+$  includes  $F$  but also other FDs such as  $A \rightarrow C$ ,  $B \rightarrow A$ , and  $C \rightarrow B$  (the latter three FDs are obtained by transitivity)
- ❑ This means that  $B \rightarrow A \in F_{R_1}$  and  $C \rightarrow B \in F_{R_2}$
- ❑ In summary,  $A \rightarrow B, B \rightarrow C, B \rightarrow A, C \rightarrow B \in F_{R_1} \cup F_{R_2}$  holds
- ❑ Consequently,  $C \rightarrow A \in (F_{R_1} \cup F_{R_2})^+$  holds due to the transitivity axiom applied to  $C \rightarrow B$  and  $B \rightarrow A$
- ❑ Hence, the decomposition of  $R$  into  $R_1(A, B)$  and  $R_2(B, C)$  with  $F_{R_1} = \{A \rightarrow B, B \rightarrow A\}$  and  $F_{R_2} = \{B \rightarrow C, C \rightarrow B\}$  preserves the FD  $C \rightarrow A$

## Dependency Preservation (IV)

- Algorithm 1 for testing dependency preservation

*bool IsDependencyPreserving1*( $\{R_1, R_2, \dots, R_n\}, F$ )

// Input: (1) A decomposition of  $R$  into the relation schemas  $R_1, R_2, \dots, R_n$

// (2) A set  $F$  of FDs on  $R$

// Output: *true*, if the decomposition is dependency preserving under  $F$ ;

// *false*, otherwise

// Step 1: Compute the closure of  $F$

$F^+ := \text{CalculateFDClosure}(F)$

// Step 2: Compute the restrictions of  $F^+$  to the  $R_i$

**for each**  $i$  in  $1..n$  **do**

$F_{R_i} := \emptyset$

**for each**  $X \rightarrow Y \in F^+$  **do**

**if**  $X \cup Y \subseteq R_i$  **then**  $F_{R_i} := F_{R_i} \cup \{X \rightarrow Y\}$

## Dependency Preservation (V)

□ Algorithm 1 for testing dependency preservation (*continued*)

// Step 3: Form the union of all restrictions

$F_r := \emptyset$

**for each**  $i$  in  $1..n$  **do**

$F_r := F_r \cup F_{R_i}$

// Step 4: Compute the closure of  $F_r$

$F_r^+ := \text{CalculateFDClosure}(F_r)$

// Step 5: Check if the two closures are equal

**return**  $(F_r^+ = F^+)$



## Dependency Preservation (VI)

- ❑ Algorithm 2 for testing dependency preservation without computing  $F^+$

*bool IsDependencyPreserving2*( $\{R_1, R_2, \dots, R_n\}, F$ )

**for each**  $X \rightarrow Y \in F$  **do**

*Result* :=  $X$

**repeat**

*OldResult* := *Result*

        // Compute the attribute closure of *Result* under  $F_r$

**for each**  $i$  in  $1..n$  **do**

            // Compute the attribute closure of *Result* under  $F_{R_i}$

$C := \text{CalculateAttributeClosure}(F, \text{Result} \cap R_i) \cap R_i$

*Result* :=  $\text{Result} \cup C$

**until** *OldResult* = *Result*

**if**  $Y \cap \text{Result} \neq Y$  **then return false**   // FD  $X \rightarrow Y$  is not preserved

**return true**   // All FDs in  $F$  are preserved

## Dependency Preservation (VII)

### □ Ideas behind Algorithm 2

- ❖ Test each FD  $X \rightarrow Y$  in  $F$  to see if it is preserved in  $F_r$  (as the union of all restrictions  $F_{R_i}$ )
- ❖ For this purpose, compute the attribute closure of  $X$  under  $F_r$ , and check whether it includes  $Y$ 
  - This is done without first computing  $F_r$  explicitly since this is quite expensive
  - The statement  $\text{CalculateAttributeClosure}(F, \text{Result} \cap R_i) \cap R_i$  computes the attribute closure of  $\text{Result}$  under  $F_{R_i}$
  - Reasons:
    - For any  $A \subseteq R_i$ ,  $A \rightarrow A^+ \in F^+$ , and  $A \rightarrow A^+ \cap R_i \in F_{R_i}$
    - Conversely, if  $A \rightarrow B \in F_{R_i}$  holds, then  $B \subseteq A^+ \cap R_i$
- ❖ The decomposition is dependency-preserving if all FDs in  $F$  are preserved
- ❖ Algorithm has a polynomial runtime complexity

## Universal Relation Assumption (I)

- ❑ The transformation of an E-R diagram into a set of relation schemas already represents an anticipated decomposition of the database schema
- ❑ But we have learned that checking whether each individual relation schema of a given decomposition satisfies a desired normal form does not guarantee a good database design
- ❑ The anticipated decomposition itself could already have a problem
- ❑ A **mandatory** requirement is lossless join property of the decomposition
- ❑ An **optional** requirement is dependency preservation of the decomposition
- ❑ In order to avoid problems and a bad database design, we have to make the **universal relation assumption**: Each normalization algorithm starts from a *single universal relation schema*  $R(A_1 : D_1, A_2 : D_2, \dots, A_n : D_n)$ , which includes *all* the attributes of the database schema, and the set  $F$  of FDs on  $R$
- ❑ The universal relation assumption contributes to the correctness criteria

## Universal Relation Assumption (II)

- ❑ This means for the relation schemas  $R_1, R_2, \dots, R_n$  obtained as the result of the transformation of an E-R diagram into relation schemas:
  - ❖ Take back the decomposition
  - ❖ If there are attributes in different relation schemas with the *same* name, make them *unique* by renaming
  - ❖ Merge the relation schemas  $R_1, R_2, \dots, R_n$  into the universal relation  $R$ , i.e.,  $R = \bigcup_{i=1}^n R_i$
- ❑ This does not make the E-R modeling process redundant since
  - ❖ the E-R diagram allows us to obtain an overview of the relevant data (i.e., entities, relationships, attributes) that have to be stored later in the database
  - ❖ each relation schema  $R_i$  is “represented” by the FD  $K_i \rightarrow R_i$  in  $F$  if  $K_i$  is the primary key of  $R_i$
- ❑ Additional FDs can and will lead to a different decomposition of  $R$