

Referential Integrity (V)

- ❑ Standard behavior of the DBMS when violating a referential integrity constraint: rejection of the action that caused the violation, i.e., the transaction performing the update action is rolled back
- ❑ Attempt to update or delete a tuple in the referenced relation R that has a matching tuple in the referencing relation S depends on the **referential action** specified using the **on delete** or **on update** sub-clauses of the **foreign key** clause
- ❑ Possible referential actions for the **on delete** sub-clause
 - ❖ **on delete cascade**: The tuple (including its primary key) of the referenced relation R as well as all the matching tuples (including their foreign keys) in a referencing relation S are deleted.
 - ❖ **on delete set null**: The tuple of the referenced relation R is deleted, and the foreign key values in all matching tuples of a referencing relation S are set to *null*. This option is only valid if the foreign key attributes in S do not have the *not null* constraint specified.

Referential Integrity (VI)

- ❑ Possible referential actions for the **on delete** sub-clause (*continued*)
 - ❖ **on delete set default**: The tuple of the referenced relation R is deleted, and the foreign key values in all matching tuples of a referencing relation S are set to their specified default values. This option is only valid if all foreign key attributes in S have a default value specified.
 - ❖ **on delete no action**: If a tuple of a relation S exists that references a tuple of a relation R , the deletion of the tuple in R is rejected. This is the default setting if the **on delete** rule is omitted.
- ❑ Possible referential actions for the **on update** sub-clause
 - ❖ **on update cascade**: A change of the value of the primary key attributes in a tuple of the referenced relation R is propagated to all tuples in a referencing relation S with the same value in their foreign key attributes.
 - ❖ **on update set null**: The value of the primary key attributes in the tuple of the referenced relation R is changed. The foreign key values in all matching tuples of a referencing relation S are set to *null*. This option is only valid if the foreign key attributes in S do not have the *not null* constraint specified.

Referential Integrity (VII)

- ❑ Possible referential actions for the **on update** sub-clause (*continued*)
 - ❖ **on update set default**: The value of the primary key attributes in the tuple of the referenced relation R is changed. The foreign key values in all matching tuples of a referencing relation S are set to their specified default values. This option is only valid if all foreign key attributes in S have a default value specified.
 - ❖ **on update no action**: If a tuple of a relation S exists that references a tuple of a relation R , the update of the primary key attribute values of the tuple in R is rejected. This is the default setting if the **on update** rule is omitted.
- ❑ Examples
 - ❖ **create table** lectures
(...,
 foreign key(held_by) **references** professors(pers-id)
 on delete cascade
 on update cascade);

Referential Integrity (VIII)

□ Examples (*continued*)

- The deletion of a tuple with the key *pers-id* in *professors* is permitted and has the effect on *lectures* that there all referencing tuples are deleted. This enables the realization of a dependent relationship.
- A change of a value of the attribute *pers-id* in *professors* is propagated to all matching values of the foreign key attribute *held_by* in *lectures*

❖ **create table** lectures

(...,

foreign key(held_by) **references** professors(pers-id)
on delete set null
on update set null);

- The value of the attribute *held_by* in *lectures* is set to *null*, if the referenced tuple with the key *pers-id* in *professors* is deleted
- The value of the attribute *held_by* of all referencing tuples in *lectures* is set to *null*, if the value of the primary key attribute *pers-id* of the referenced tuple in *professors* is changed

General Integrity

- ❑ An **assertion** is a predicate that expresses a condition which is to be always satisfied by a database system
- ❑ Domain constraints and referential ICs are special kinds of assertions
- ❑ There are also conditions which cannot be expressed with these two kinds like conditions with respect to several relations
- ❑ Syntax: **create assertion** <assertion name> **check** <condition>
- ❑ Example: There must be at least four professors in order to maintain teaching
create assertion AlwaysFourProfessors
check (4 <= (**select count**(*) **from** professors))
- ❑ The DBMS tests an assertion for validity first, and if the assertion is valid, future modifications of the database are only allowed if the assertion is not violated
- ❑ Complex assertions can lead to an overhead; they should be used with care
- ❑ Most DBS have not implemented assertions

Triggers (I)

❑ A trigger

- ❖ is a user-defined procedure that the DBMS executes automatically if a certain condition is fulfilled or as a side effect of a modification of the database
- ❖ is a general and powerful mechanism for maintaining data(base) consistency, monitoring database updates, and updating database statistics
- ❖ cannot only deploy check functions but also computation functions

❑ A trigger is based on the ECA (Event-Condition-Action) model and has three components:

- ❖ The event(s) to which it reacts
 - These are usually database manipulation operations (insert, delete, update) that are explicitly applied to the database
 - Person who writes the trigger must ensure that all possible events are accounted for
 - Trigger must specify if it should be executed *before* or *after* an event

Triggers (II)

- ❖ The **condition** that determines whether the action should be executed
 - If no condition is specified, the action will be executed once the event occurs
 - If a condition is specified in the *when* clause of the trigger, it is first evaluated, and only if it evaluates to *true*, the action will be executed
- ❖ The **action(s)** to be taken
 - The action is usually a sequence of SQL statements
 - It can also be a database transaction or an external program that will be automatically executed

❑ Part of the SQL syntax for creating a trigger in Oracle

- ❖ **create trigger** <trigger name>
 {**before** | **after**}
 {**insert** | **delete** | **update** [**of** <column list>]} **on** <table name>
 [**for each row**]
 [**when** <condition>]
 <trigger actions>;

Triggers (III)

- ❑ In Oracle the trigger actions are implemented in PL/SQL, which is Oracle's database application programming language
- ❑ Example: Trigger preventing that professors can be demoted by a rank

```
create trigger noDemotion  
before update on professors  
for each row  
when (:old.rank is not null)  
begin  
    if :old.rank = "C3" and :new.rank = "C2" then :new.rank = "C3" end if;  
    if :old.rank = "C4" then :new.rank = "C4" end if;  
    if :new.rank is null then :new.rank = :old.rank end if;  
end
```

- ❑ The clause **create trigger** provides a trigger specification
- ❑ The time of releasing the trigger body can be **before** or **after** the data manipulation operation that has released the trigger

Triggers (IV)

❑ Trigger events

- ❖ **update** [**of** <column1, column2, ...>] **on** <relation name>
- ❖ **insert on** <relation name>
- ❖ **delete on** <relation name>

❑ A trigger can be defined for one or several events. In case of several events, a case distinction can be expressed in the body through the clauses

- ❖ **if updating** [<column1, column2, ...>] **then** ...
- ❖ **if inserting then** ...
- ❖ **if deleting then** ...

❑ Trigger type

- ❖ A **statement-level trigger** (default) is released exactly *once* either before or after the respective event.
- ❖ A **row-level trigger** is indicated by the **for each row** clause and signifies that the trigger is executed *for each changed tuple* in the table

Triggers (V)

- ❑ Distinguishing tuples before and after a change in a row-oriented trigger
 - ❖ Keyword **:old** refers to a deleted tuple or to a tuple before it was updated
 - ❖ Keyword **:new** refers to a newly inserted or newly updated tuple
 - ❖ This is the only possibility to address tuples before or after a change
 - ❖ Another access to the relation is not possible any more, even if the relation name would be used in the respective block
- ❑ Trigger restriction
 - ❖ The **when** predicate specifies a condition that must be satisfied to release the execution of the trigger body
 - ❖ If a row-oriented trigger is used, the new resp. old tuple of the relation can be addressed by the keywords **:new** resp. **:old**
- ❑ Trigger body / trigger actions
 - ❖ SQL queries
 - ❖ Oracle-specific PL/SQL commands with the aforementioned extensions

Triggers (VI)

- ❑ Example: Protocol of the changes of the attribute *salary* of a relation *Persons*

```
create trigger StoreSalary  
before update on Persons  
for each row  
when (:old.salary > 1500)  
begin insert into diff values (:old.salary, :new.salary, sysdate) end;
```

- ❑ Example: Check at insertion time that a salary increase is inapplicable to persons with a monthly salary greater than \$10,000

```
create trigger CheckSalary  
before update on Persons  
for each row when (:new.salary > 10000)  
begin  
    :new.salary := :old.salary;    // assignment only possible for before update  
end;
```