

Database Management Systems (COP 5725)
(Spring 2009)

Instructor: Dr. Markus Schneider

TA:
Young Namkoong
Xiao Li

Exam 2 Solution

Name:	
UFID:	
Email Address:	

Pledge (Must be signed according to UF Honor Code)

On my honor, I have neither given nor received unauthorized aid in doing this assignment.

_____Signature

For scoring use only:

	Maximum	Received
Exercise 1	20	
Exercise 2	37	
Exercise 3	25	
Exercise 4	18	
Total	100	

Exercise 1 (SQL-I)

[20 points]

A) Consider the following database schema: (6 points)

```
EMP (EMP_ID, Dept_num)
```

Assume that we have four queries that operate on the above schema:

- I. (SELECT * FROM EMP WHERE Emp.Dept_num = 100)
 INTERSECT
 (SELECT * FROM EMP WHERE Emp.Dept_num = 200)
- II. SELECT * FROM EMP
 WHERE Emp.Dept_num = 100 AND Emp.Dept_num = 200
- III. (SELECT * FROM EMP WHERE Emp.Dept_num = 100)
 UNION
 (SELECT * FROM EMP WHERE Emp.Dept_num = 200)
- IV. SELECT * FROM EMP
 WHERE Emp.Dept_num = 100 OR Emp.Dept_num = 200

- 1) List all (if any) of the above queries that correctly give the answers if we wanted to find all people working in department #100 *and* department #200: (3 points)

I

- 2) List all (if any) of the above queries that correctly give the answers if we wanted to find all people working in department #100 *or* department #200: (3 points)

III, IV

B) Consider the following database schema: **(14 points)**

```
INVESTORS (Name, Phone, Address, Country)
OWNS (StockID, InvestorName)
STOCKS (Id, CompanyName, Country)
```

INVESTOR contains a tuple for each investor, recording his/her name, phone, address, and country of residence.

STOCKS records the identifier and company name of each stock and the country where the company is located. For example, the stock identifier of the company “Sun Microsystems Inc.” can be “JAVA” (similar to the NASDAQ ticker symbol) and the country will be “United States”.

1) Write the SQL statement for the following query. (2 points)

Retrieve the name of each investor who owns the stock of the company named “Oracle”.

```
SELECT O.InvestorName
FROM OWNS O, STOCKS S
WHERE O.StockID = S.Id AND S.CompanyName = 'Oracle'
```

2) Write the SQL statement for the following query. (3 points)

Retrieve the name of each investor who owns both the stock of the company with the stock id “SYBS” (Sybase) and the company with the stock id “IBM”.

```
SELECT O1.InvestorName
FROM OWNS O1, OWNS O2
WHERE O1.StockID = 'SYBS' AND O2.StockID = 'IBM' AND
      O1.InvestorName = O2.InvestorName
```

3) Write the SQL statement for the following query. (3 points)

Retrieve the name of each investor who owns the stock of the company which is located in the same country as the investor himself/herself.

```
SELECT O.InvestorName
FROM OWNS O, INVESTORS I, STOCKS S
WHERE O.StockID = S.Id AND O.InvestorName = I.Name AND
      S.Country = I.Country
```

- 4) Write the SQL statement for the following query. (3 points)

For each stock, retrieve its identifier, name and the total number of investors that own this stock.

```
SELECT S.Id, S.CompanyName, COUNT(*)  
FROM OWNS O, STOCKS S  
WHERE S.ID = O.StockId  
GROUP BY S.Id, S.CompanyName
```

- 5) Describe the output of the following SQL statement in one English statement. (3 points)

```
SELECT I.Name  
From INVESTORS I  
WHERE NOT EXISTS  
    (SELECT S.ID  
     FROM STOCKS S  
     WHERE S.Country = 'US' AND  
           NOT EXISTS (SELECT O.StockID  
                      FROM OWNS O  
                      WHERE O.StockID = S.Id AND  
                            O.InvestorName = I.Name));
```

It retrieves the name of the investors who own the stock of all American companies.

Exercise 2 (SQL-II)

[37 points]

A) Consider the following database schema: **(9 points)**

COPLIST (X, Y, Z)

Assume that we have the following table:

COPLIST

<u>X</u>	Y	Z
1	ATC-LaLaLa	20
2	Barbie_Girl	10
3	Barbie_Girl	10
4	Barbie_Girl	30
5	Come_Clean	40

Answer the query result to the following SQL statements.

1) SELECT COUNT(X) FROM COPLIST; (1 point)

5

2) SELECT COUNT(DISTINCT X) FROM COPLIST; (2 points)

5

3) SELECT COUNT(Y) FROM COPLIST; (1 point)

5

4) SELECT COUNT(DISTINCT Y) FROM COPLIST; (2 points)

3

5) SELECT COUNT(Z) FROM COPLIST; (1 point)

5

6) SELECT COUNT(DISTINCT Z) FROM COPLIST; (2 points)

4

B) Consider the following database schema: (8 points)

Student (UF_ID, Name, Dept, GPA)

Assume that we have the following table:

Student			
UF_ID	Name	Dept	GPA
1111-2345	Martin Fitzgerald	CISE	3.70
3439-3984	Elena Delgado	ECE	3.60
1847-9911	Shirley Schmidt	ECE	2.90
7637-2521	Alan Shore	CISE	3.30
1199-0174	Denny Crane	CISE	3.50
3976-6158	Tony Almeida	ECE	3.10
3342-2850	Dr. Ross Geller	ECE	3.40

1) Show the results after the following query is executed: (2 points)

```
SELECT UF_ID, Name
FROM Student
WHERE Dept = 'CISE';
```

```
UF_ID      Name
-----
1111-2345  Martin Fitzgerald
7637-2521  Alan Shore
1199-0174  Denny Crane
```

2) Show the results after the following query is executed: (3 points)

```
SELECT UF_ID, Name
FROM Student
WHERE GPA > (SELECT AVG(GPA) FROM Student);
```

```
UF_ID      Name
-----
1111-2345  Martin Fitzgerald
3439-3984  Elena Delgado
1199-0174  Denny Crane
3342-2850  Dr. Ross Geller
```

3) Show the results after the following query is executed: (3 points)

```
SELECT Dept, AVG(GPA)
FROM Student
GROUP BY Dept
HAVING AVG(GPA) > 3.3;
```

```
Dept      AvgGPA
-----
CISE      3.5
```

C) Consider the following database schema: **(20 points)**

Department		
Name	Position	Salary
Socrates	TA	100.00
Bill Gates	TA	130.00
Steve Jobs	Prof	200,000.00
Larry Page	Prof	500,000.00
Arkimedes	Prof	350,000.00

On this relation, we will create a view called high_salary as below:

```
CREATE VIEW high_salary AS
SELECT Name, Position, Salary
FROM Department
WHERE Salary > 200.00
```

Assume that our DBMS allow insertion into views.

Based upon the above assumption, let's execute the following command:

```
INSERT INTO Department VALUES('Jay Leno', 'Chairman', 120,000.00);
```

- 1) Show the results of the following SQL command **after the above insertion.**
(3 points)

```
SELECT * FROM high_salary;
```

Name	Position	Salary

Steve Jobs	Prof	200,000.00
Larry Page	Prof	500,000.00
Arkimedes	Prof	350,000.00
Jay Leno	Chairman	120,000.00

- 2) Show the results of the following SQL command **after the above insertion.**
(3 points)

```
SELECT * FROM Department;
```

Name	Position	Salary

Socrates	TA	100.00
Bill Gates	TA	130.00
Steve Jobs	Prof	200,000.00
Larry Page	Prof	500,000.00
Arkimedes	Prof	350,000.00
Jay Leno	Chairman	120,000.00

- 3) What would be the answer of (1) if we had inserted the Chairman Jay Leno tuple in the relation high_salary instead? (3 points)

Name	Position	Salary

Steve Jobs	Prof	200,000.00
Larry Page	Prof	500,000.00
Arkimedes	Prof	350,000.00
Jay Leno	Chairman	120,000.00

- 4) Will the following command work? Why or why not? (assume that this takes place after the above insertion into Department). (3 points)

```
INSERT INTO high_salary VALUES('George Clooney', 'TA', 80.00);
```

It can work.

We can insert tuples into a view even though new data might not satisfy this view's constraint.

- 5) Show the results of the following SQL command **after the two insertion commands above**. (i.e., shown at the problem statements and in Exercise2-(C)-(4)). (4 points)

```
SELECT * FROM high_salary;
```

Name	Position	Salary

Steve Jobs	Prof	200,000.00
Larry Page	Prof	500,000.00
Arkimedes	Prof	350,000.00
Jay Leno	Chairman	120,000.00

- 6) Show the results of the following SQL command **after the two insertion commands above**. (i.e., shown at the problem statements and in Exercise2-(C)-(4)). (4 points)

```
SELECT * FROM Department;
```

Name	Position	Salary

Socrates	TA	100.00
Bill Gates	TA	130.00
Steve Jobs	Prof	200,000.00
Larry Page	Prof	500,000.00
Arkimedes	Prof	350,000.00
George Clooney	TA	80.00

Exercise 3 (Relational Algebra)

[25 points]

Consider the following schema for a college database:

DEPARTMENT (dept_id, dept_name, dept_chairman)
PROFESSOR (prof_name, dept_id, building, office, email)
STUDENT (student_id, student_name, start_year, dept_id)
EXAM (student_id, course_id, prof_name, grade)

Express the following queries in relational algebra expression (without using the aggregate functions such as AVG, MAX, ALL, and COUNT, etc.)

(A) Find the IDs of all students that started in 2006 or later. (3 points)

$$\pi_{\text{student_id}} (\sigma_{\text{start_year} \geq 2006} (\text{STUDENT}))$$

(B) Find the IDs and names of all students in the CISE department that started in 2007. (5 points)

$$\begin{aligned} R1 &\leftarrow (\text{STUDENT} \bowtie_{\text{STUDENT.dept_id}=\text{DEPARTMENT.dept_id}} (\text{DEPARTMENT})) \\ \pi_{\text{student_id}, \text{student_name}} (\sigma_{\text{start_year}=2007 \wedge \text{dept_name}='CISE'} (R1)) \end{aligned}$$

(C) Find the name of the professor(s) who manage exams both for 'COP5725' and 'CIS4301'. (5 points)

$$\pi_{\text{prof_name}} (\sigma_{\text{course_id}='COP5725'} (\text{EXAM})) \cap \pi_{\text{prof_name}} (\sigma_{\text{course_id}='COP4301'} (\text{EXAM}))$$

(D) Find the IDs of the students who have had an exam given by a professor who is currently chairman of a department. (5 points)

$$\pi_{\text{student_id}} (\text{EXAM} \bowtie_{\text{prof_name}=\text{dept_chairman}} \text{DEPARTMENT})$$

(E) Find the IDs and names of all students who got grade A in all their exams. (7 points)

$$\begin{aligned} &\rho(R1, \pi_{\text{student_id}} (\text{EXAM}) - \pi_{\text{student_id}} (\sigma_{\text{grade} \neq 'A'} (\text{EXAM}))) \\ &\pi_{\text{student_id}, \text{student_name}} (\text{STUDENT} \bowtie R1) \end{aligned}$$

Exercise 4 (Relation Calculus and QBE)

[18 points]

Consider the following schema for an airline database:

FLIGHTS (flight_num, from_city, to_city, distance, depart_time, arrival_time)

AIRCRAFT (aircraft_id, aircraft_name, cruising_range)

PILOTS (pilot_id, pilot_name, salary)

CERTIFIED (pilot_id, aircraft_id)

(A) Write the *tuple relational calculus expression* to the following query. (4 points)

“Find the pilot_ids of pilots certified for ‘AirTran’ aircraft.”

$$\{C.pilot_id \mid C \in \text{CERTIFIED} \wedge \\ \exists A \in \text{AIRCRAFT}(A.aircraft_id = C.aircraft_id \wedge A.aircraft_name = 'AirTran')\}$$

(B) Write the *domain relational calculus expression* to the same query in problem (A). (4 points)

$$\{\langle Cpid \rangle \mid \langle Cpid, Caid \rangle \in \text{CERTIFIED} \wedge \exists Aid, Aname, Arange \\ (\langle Aid, Aname, Arange \rangle \in \text{AIRCRAFT} \wedge Aid = Caid \wedge Aname = 'AirTran')\}$$

(C) Write the *tuple relational calculus expression* to the following query. (5 points)

“Find the aircraft_ids of all aircrafts that can be used on non-stop flights from ‘Gainesville’ to ‘Dortmund’.”

$$\{A.aircraft_id \mid A \in \text{AIRCRAFT} \wedge \exists F \in \text{FLIGHTS} \\ (F.from_city = 'Gainesville' \wedge F.to_city = 'Dortmund' \wedge A.cruising_range > F.distance)\}$$

(D) Briefly explain QBE (Query-by-Example). (5 points)

1. Query-by-Example (QBE) is a database query language for relational database.
2. QBE is developed by IBM, later part of DB2.
3. QBE is based on the domain relational calculus: variables are bound to attribute domains
4. QBE queries are expressed by means of examples. The system generalizes the examples in order to compute answers for queries.